

In []:

EDA based on Heart Disease

```
In [64]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from scipy import stats
```

```
In [65]: df=pd.read_csv('heart.csv')
df.head()
```

```
Out[65]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2



```
In [66]: df.shape
```

```
Out[66]: (1025, 14)
```

```
In [67]: df.dtypes
```

```
Out[67]: age          int64
sex            int64
cp             int64
trestbps       int64
chol           int64
fbs            int64
restecg        int64
thalach        int64
exang          int64
oldpeak        float64
slope          int64
ca             int64
thal           int64
target         int64
dtype: object
```

```
In [68]: print(df.columns.tolist())
```

```
['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach', 'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target']
```

```
In [69]: print(df.info())
```

```

<class 'pandas.DataFrame'>
RangeIndex: 1025 entries, 0 to 1024
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         1025 non-null   int64
 1   sex         1025 non-null   int64
 2   cp          1025 non-null   int64
 3   trestbps    1025 non-null   int64
 4   chol        1025 non-null   int64
 5   fbs         1025 non-null   int64
 6   restecg     1025 non-null   int64
 7   thalach     1025 non-null   int64
 8   exang       1025 non-null   int64
 9   oldpeak     1025 non-null   float64
10   slope       1025 non-null   int64
11   ca          1025 non-null   int64
12   thal        1025 non-null   int64
13   target      1025 non-null   int64
dtypes: float64(1), int64(13)
memory usage: 112.2 KB
None

```

Before performing exploratory data analysis, We checked the dataset for missing values.

```

In [70]: #checking for null values
df.isnull().sum()

```

```

Out[70]: age         0
sex         0
cp          0
trestbps    0
chol        0
fbs         0
restecg     0
thalach     0
exang       0
oldpeak     0
slope       0
ca          0
thal        0
target      0
dtype: int64

```

The result showed zero missing values across all 14 variables, so no missing-value treatment was required.

```

In [71]: #checking duplicates
print("\n The no.of duplicate rows:", df.duplicated().sum())

```

The no.of duplicate rows: 723

```

In [72]: print("Number of unique rows:", df.drop_duplicates().shape[0])

```

Number of unique rows: 302

```

In [73]: #removing duplicates
df_clean=df.drop_duplicates().copy()
df_clean.head()

```

Out[73]:

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal
0	52	1	0	125	212	0	1	168	0	1.0	2	2	3
1	53	1	0	140	203	1	0	155	1	3.1	0	0	3
2	70	1	0	145	174	0	1	125	1	2.6	0	0	3
3	61	1	0	148	203	0	1	161	0	0.0	2	1	3
4	62	0	0	138	294	1	1	106	0	1.9	1	3	2



In [74]: `print(df_clean.shape)`

(302, 14)

A major data-cleaning step was duplicate detection. Out of 1,025 records, 723 were duplicate rows. After removing these duplicates, the dataset contained 302 unique patient records. This cleaned dataset was then used for further analysis.

In [75]: `df_clean.describe().T`

Out[75]:

	count	mean	std	min	25%	50%	75%	max
age	302.0	54.420530	9.047970	29.0	48.00	55.5	61.00	77.0
sex	302.0	0.682119	0.466426	0.0	0.00	1.0	1.00	1.0
cp	302.0	0.963576	1.032044	0.0	0.00	1.0	2.00	3.0
trestbps	302.0	131.602649	17.563394	94.0	120.00	130.0	140.00	200.0
chol	302.0	246.500000	51.753489	126.0	211.00	240.5	274.75	564.0
fbs	302.0	0.149007	0.356686	0.0	0.00	0.0	0.00	1.0
restecg	302.0	0.526490	0.526027	0.0	0.00	1.0	1.00	2.0
thalach	302.0	149.569536	22.903527	71.0	133.25	152.5	166.00	202.0
exang	302.0	0.327815	0.470196	0.0	0.00	0.0	1.00	1.0
oldpeak	302.0	1.043046	1.161452	0.0	0.00	0.8	1.60	6.2
slope	302.0	1.397351	0.616274	0.0	1.00	1.0	2.00	2.0
ca	302.0	0.718543	1.006748	0.0	0.00	0.0	1.00	4.0
thal	302.0	2.314570	0.613026	0.0	2.00	2.0	3.00	3.0
target	302.0	0.543046	0.498970	0.0	0.00	1.0	1.00	1.0

The descriptive statistics show that the average age of patients is around 54 years. The average resting blood pressure is approximately 132, while the average cholesterol level is around 247. The average maximum heart rate is approximately 150 beats per minute.

In [76]: `percentage_sex=df_clean['sex'].value_counts(normalize=True)*100`
`print(percentage)`

```
sex
1    68.211921
0    31.788079
Name: proportion, dtype: float64
```

```
In [77]: percentage_exang=df_clean['exang'].value_counts(normalize=True)*100
print(percentage_exang)
```

```
exang
0    67.218543
1    32.781457
Name: proportion, dtype: float64
```

```
In [78]: percentage_target=df_clean['target'].value_counts(normalize=True)*100
print(percentage_target)
```

```
target
1    54.304636
0    45.695364
Name: proportion, dtype: float64
```

The average age of patients is 54.42 years, with a mean cholesterol level of 246.5 mg/dL and a mean resting blood pressure of 131.6 mmHg. The average maximum heart rate achieved is 149.6 bpm. Approximately 68% of patients are male, 33% experienced exercise-induced angina(chest pain), and 54% have heart disease.


```
In [79]: #Target variable analysis
df_clean['target'].value_counts()
```

```
Out[79]: target
1    164
0    138
Name: count, dtype: int64
```

```
In [80]: # 1.Bar Chart – Heart Disease Outcome

sns.countplot(data=df_clean, x='target')
plt.title('Heart Disease Distribution')
plt.xlabel('Target Variable (0 = No Disease, 1 = Disease)')
plt.ylabel('No. of Patients')

plt.ylim(0, 200)
plt.show()
```




The target variable represents whether a patient has heart disease. After cleaning, 164 patients were classified as having heart disease, while 138 were classified as not having heart disease. The distribution is relatively balanced, although the disease group is slightly larger.

```
In [81]: # 2. Histogram – Age Distribution

plt.hist(df_clean['age'], bins=15, edgecolor='black')
plt.title('Age Distribution of Patients')
plt.xlabel('Age')
plt.ylabel('No. of Patients')
plt.ylim(0, 50)

plt.show()
```




The age histogram shows that the dataset contains patients mainly in the middle-aged and older age groups. The overall age range is 29 to 77 years, with an average age of approximately 54 years. Most patients are middle-aged or older, making age an important variable to investigate in relation to heart disease.

```
In [82]: # 3. Boxplot

sns.boxplot(data=df_clean, x='target', y='age')

plt.title('Age vs Heart Disease')
plt.xlabel('Heart Disease (0 = No Disease, 1 = Disease)')
plt.ylabel('Age')
plt.show()
```




This boxplot compares the age distribution between the two target groups. By comparing the median and spread of the two groups, we can observe whether patients with heart disease tend to be older or younger than patients without disease.

```
In [83]: sns.boxplot(data=df_clean, x='target', y='thalach')

plt.title('Heart Disease vs Maximum Heart Rate')
plt.xlabel('Heart Disease (0 = No Disease, 1 = Disease)')
plt.ylabel('Maximum Heart Rate')


plt.show()
```



This boxplot compares the maximum heart rate achieved by patients with and without heart disease. The distributions show that maximum heart rate differs between the two groups. We can see that in the above box plot, heart disease patients have a higher median maximum heart rate (~160) and a wider spread compared to no heart disease patients, who have a lower median (~142) with a few outliers. This suggests that maximum heart rate is a potentially useful feature for classifying between the two classes in this dataset.

```
In [84]: plt.figure(figsize=(8, 5))

sns.countplot(data=df_clean, x='cp', hue='target')
plt.title('Chest Pain Type vs Heart Disease')
plt.xlabel('Chest Pain Type')
plt.ylabel('No. of Patients')
plt.legend(title='Heart Disease', labels=['No Disease', 'Disease'])
plt.show()
```



This graph shows the relationship between chest pain and heart disease. It compares patients with and without heart disease for each type of chest pain. We can see which chest pain types are more common in heart-disease patients.


```
In [85]: #5. Correlation Heatmap

plt.figure(figsize=(12, 8))

corr_matrix = df_clean.corr(numeric_only=True)

sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5,
plt.title('Correlation Heatmap of Heart Disease')

plt.show()
```




The heatmap helps identify potentially important relationships and possible multicollinearity between features, but correlation alone does not establish a causal relationship.

```
In [86]: #6. Scatter Plot – Age vs Maximum Heart Rate

sns.scatterplot(data=df_clean, x='age', y='thalach', hue=df_clean['target'].map(

plt.title('Age vs. Maximum Heart Rate')
plt.xlabel('Age')
plt.ylabel('Maximum Heart Rate')

plt.legend(title='Heart Disease', )
plt.show()
```



The scatter plot helps us visualize how maximum heart rate changes with age and whether the pattern differs between patients with and without heart disease. Generally, maximum heart rate tends to vary with age, and the target-based grouping allows us to observe differences between the two patient groups.

In []: